

**Master in Advanced Mathematics and Mathematical Engineering**

# **Transparent Live Migration of Container Deployments in Userspace**

**Master's Thesis - Oral Defense**

Carlos Segarra - `carlos.segarra@estudiant.upc.edu`

Advisor: Jordi Guitart - `jguitart@ac.upc.edu`

Barcelona, Tuesday July 7, 2020



UNIVERSITAT POLITÈCNICA DE CATALUNYA  
BARCELONATECH

Facultat de Matemàtiques i Estadística

UNIVERSITAT POLITÈCNICA DE CATALUNYA  
BARCELONATECH

Facultat d'Informàtica de Barcelona



# Session Outline

## 1 Introduction

- Motivation & Main Goal
- Contributions List

## 2 Background Concepts

- Containers and runC
- Checkpoint-Restore and Live Migration
- CRIU: Checkpoint-Restore in Userspace

## 3 Implementation

- Diskless Migration
- Iterative Migration
- Established TCP Connections
- Putting it All Together

## 4 Evaluation

- Application Downtime
- Comparison with Other Tools

## 5 Conclusions & Future Work



### Problem Motivation:

- **Containers** have become the go-to alternative for sandboxing and deploying distributed applications, they build on the concept of **namespaces**.
- **Checkpointing** provides systems with **fault-tolerance** and **state consistency** enabling rollback and restore from previous stable versions.
- **Checkpoint-Restore in Userspace (CRIU)** is a tool to perform checkpoint-restore of processes, transparently to the user.
- **Live migration** consists in transparently relocating running services for improved resource-management and increased QoS.

### Main Goal

Design, implement, and evaluate a tool for efficient live migration of running runC containers using **CRIU**.



# Introduction

## Current Limitations & Contributions List

### Current limitations:

- Limited integration of C/R within container engines (none in terms of live migration).
- VMs are losing ground to containers, yet no replacement for VM migration tools.
- Existing live-migration libraries have lots of dependencies and are complex to set up.

### Main Objectives:

- Implement a fully-featured container migration tool.
- Support network and memory intensive containers.

### Contributions List

- ① An exhaustive micro-benchmark of different CRIU functionalities.
- ② An open-source library for live migration of runC containers using CRIU.
- ③ An easy-to-use binary to transparently migrate containers.
- ④ A comparison of our solution with traditional VM migration.



### Linux Containers:

#### Definition

A linux container is a set of processes isolated from the rest of the system.

- Containers rely on **namespaces** and **control groups** for fine-grained resource control.
- A **container engine** transforms a *container image* into a process by means of a *runtime*.
- We choose **runC** as our container runtime as it has the best integration with CRIU.

#### Introduction to runC:

- **runC** is the Open Container Initiative's (OCI) reference runtime implementation.
- Rather than **images**, it relies on OCI bundles: **config.json** + **rootfs**.
- It sits at the core of most container engines as depicted in Figure 1.



### Introduction to runC:

- **runC** is the Open Container Initiative's (OCI) reference runtime implementation.
- Rather than **images**, it relies on OCI bundles: **config.json** + **rootfs**.
- It sits at the core of most container engines as depicted in Figure 1.

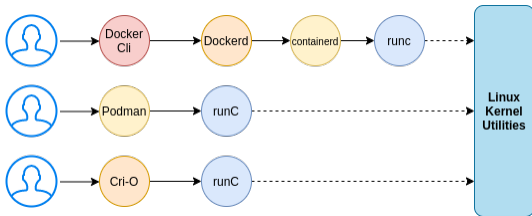


Figure 1: Placing runC in the call stack of different container engines.



# Background Concepts

## Checkpoint-Restore and Live Migration

### Checkpoint-Restore:

#### Definition (Encyclopedia of Parallel Computing)

Checkpointing refers to the ability to store the state of a computation in a way that allows it be continued at a later time without changing the computation's behavior.

- The **state** is made up of: memory, pipes, sockets, ...
- This state is used to **restore** the process.
- Originated in HPC and popularized through VMs.
- **Live migration** consists in checkpointing in an environment, and restoring in a different one, **without affecting the application's availability**.

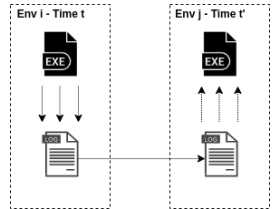


Figure 2: Basic working principle of live migration.



# Background Concepts

## CRIU: Checkpoint-Restore in Userspace

### CRIU: Checkpoint-Restore in Userspace

- We rely on **CRIU** to perform C/R from userspace, **transparently** to the user.
- To **checkpoint**, CRIU gathers information from different resources and dumps the content to files.
- To **restore**, it morphs itself into the to-be restored process.
- To achieve **efficient live migration**, several design choices must be made.

	CRIU	DMTCP	BLCR
<b>Target App.</b>	Containers	HPC	HPC
<b>Standard Kernel</b>	> 3.11	Yes	Yes
<b>Transparency</b>	Full	Pre-Load	Pre-Load
<b>Unmodified App.</b>	Yes	Yes	No
<b>Containers</b>	Yes	No	No
<b>Distributed App.</b>	No	Yes	Yes
<b>Open Files</b>	Yes	No	No

Table 1: Comparison with other popular C/R tools.





# Implementation

## Design Choices: Diskless Migration

### How to reduce the impact of file I/O to disk?

- Avoid writing files twice (on dump + transfer) → Use CRIU's **page-server**.
- Avoid writing to disk at all → Use a **tmpfs** mount rather than disk.
- We compare a small application (100 kB, left) and a big one (1 GB, right).

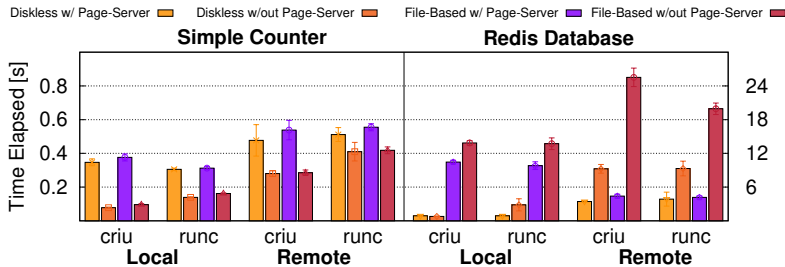


Figure 3: Diskless migration micro-benchmark.



# Implementation

## Design Choices: Iterative Migration

### How to minimize the application's downtime during migration?

- Perform iterative dumps → Use **pre-dump** and memory tracking.
- Ensure freshness upon restore → Parent directories linked list.
- For each app, we compare a setting w/out memory changes (left) and one with (right).

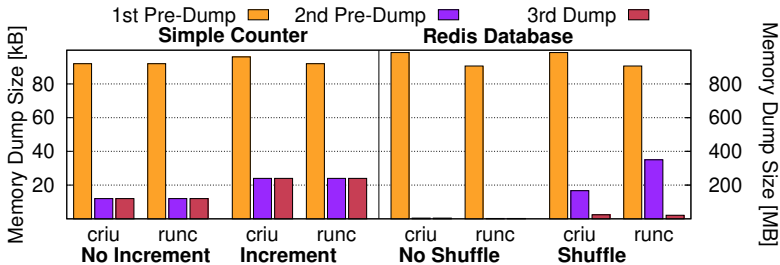


Figure 4: Iterative migration micro-benchmark.



# Implementation

## Design Choices: Established TCP Connections

### How to migrate established TCP connections?

- Rely on `TCP_REPAIR` socket option.
- Re-create `iptables` on remote end.
- Re-use same IP address:
  - ① Using network namespaces.
  - ② Using locally scoped addresses.
- In the experiments we present:
  - ① Downtime after an extended stop (top).
  - ② Reactivity to immediate C/R (bottom).

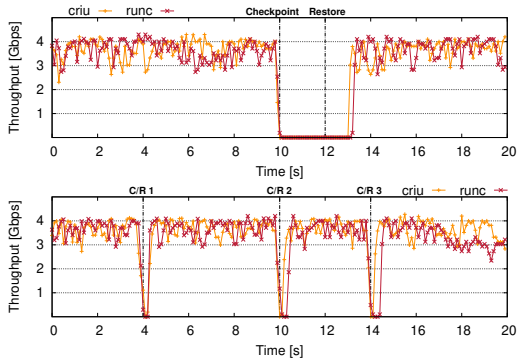


Figure 5: TCP connections micro-benchmark.



# Implementation

## Putting it all together

### Efficient live migration of running containers:

- We employ a single process (no listening daemon on destination).
- We require **CAP\_SYS\_ADMIN** capabilities.
- We make two important assumptions:
  - ① The OCI bundle is available on destination.
  - ② The user has SSH access to both machines.
- To run with default parameters we require:
  - ① A running **runC** container's name.
  - ② An IP address where to migrate it.

```
1 while size_to_xfer > MEMORY_THRESHOLD do
2     prepare_migration();
3     start_page_server();
4     pre_dump();
5     transfer_intermediate_files();
6     link_directories();
7 end while
8 start_page_server();
9 dump();
10 transfer_files();
11 restore_on_remote();
```

Algorithm 1: Main migration loop.



#### Impact of the threshold value in total downtime:

- **Downtime** (time container is unresponsive) is the key metric of success in live migration.
- The **MEMORY\_THRESHOLD** becomes a decisive design parameter.
- We study its impact on downtime when migrating a Redis in-memory DB.

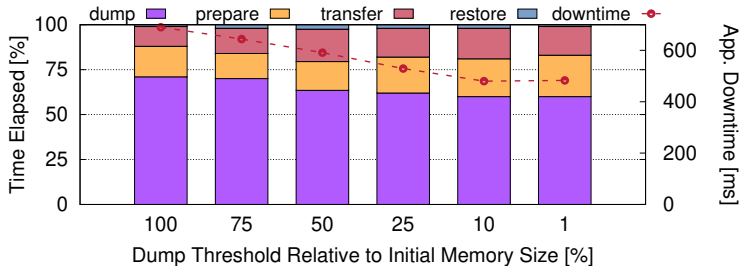


Figure 6: Application downtime macro-benchmark.



# Evaluation

## Comparison w/ Other Techniques

### Comparison with other migration techniques:

- We compare our performance with two different settings:
  - 1 Naive migration: dump-transfer-restore.
  - 2 VM live migration: using VirtualBox's teleport.
- We measure:
  - 1 Scalability regarding the container's memory size.
  - 2 Overall time elapsed (and breakdown).
- We conclude:
  - 1 Higher baseline than naive due to increased set-up.
  - 2  $\times 1.5$  slowdown vs.  $\times 2$  for VM and  $\times 10$  for naive.

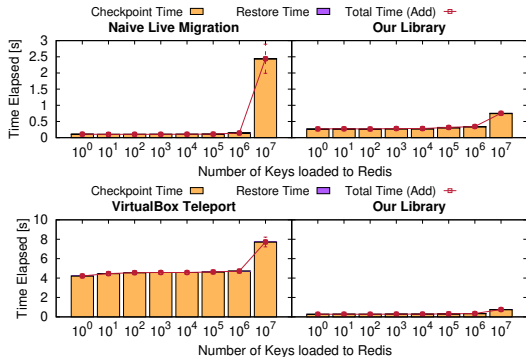


Figure 7: Comparison w/ naive migration and VM migration.



# Conclusions & Future Work

## Conclusions

- + Implemented an easy-to-use efficient solution.
- + Has minimal dependencies and requires minimal set-up.
- Limited comparison against other solutions (several dependencies and hard to set-up).
- Implementation does not support all features initially planned.

## To-Do List

- Technical Details: upon completion we plan on submitting to a specialized conference.
  - ① Evaluate against other live migration solutions for containers (e.g. **P.Haul**).
  - ② Evaluation against other VM migration tools (e.g. **KVM**'s or **LXC**'s).
  - ③ Circumvent the pre-existence requirement for OCI bundles → Generate during **pre-dump** phase.
- Future lines of work:
  - ① Initial goal was to support migration of distributed deployments.
    - ① Implement distributed checkpointing algorithms.
    - ② Integrate with container orchestrators.



Thank you for your attention,

**Observations, Doubts & Suggestions Welcome**

Follow the development:

<https://github.com/live-containers/live-migration>

[carlos.segarra@estudiant.upc.edu](mailto:carlos.segarra@estudiant.upc.edu)

